

Alive2: Bounded Translation Validation for LLVM

Nuno.P.Lopes, Juneyoung Lee, Chung-Kil Hur, Zhenyang Liu, John Regehr

Presented by JS. Kwon

IS661 Paper Review

Background: Compiler's Optimization



```
int foo(int X) {  
    int result = X % 4;  
    return result;  
}
```

High-cost Program



```
int foo(int X, int Y) {  
    int result = X & 3;  
    return result;  
}
```

Low-cost Program

Background: Compiler's Optimization



```
int foo(int X) {  
    int result = X % 4;  
    return result;  
}
```

High-cost Program



```
int foo(int X, int Y) {  
    int result = X & 3;  
    return result;  
}
```

Low-cost Program



```
void swap(int* X, int* Y) {  
    int tmp = *X;  
    *X = *Y;  
    *Y = tmp;  
    return;  
}
```

High-cost Program



```
void swap(int* X, int* Y) {  
    *X = *X ^ *Y;  
    *Y = *X ^ *Y;  
    *X = *X ^ *Y;  
    return;  
}
```


Low-cost Program

Background: Compiler's Optimization




```
int foo(int X) {  
    int result = X % 4;  
    return result;  
}
```

High-cost Program



Optimization




```
int foo(int X, int Y) {  
    int result = X & 3;  
    return result;  
}
```

Low-cost Program




```
void swap(int* X, int* Y) {  
    int tmp = *X;  
    *X = *Y;  
    *Y = tmp;  
    return;  
}
```

High-cost Program



Optimization



```
void swap(int* X, int* Y) {  
    *X = *X ^ *Y;  
    *Y = *X ^ *Y;  
    *X = *X ^ *Y;  
    return;  
}
```

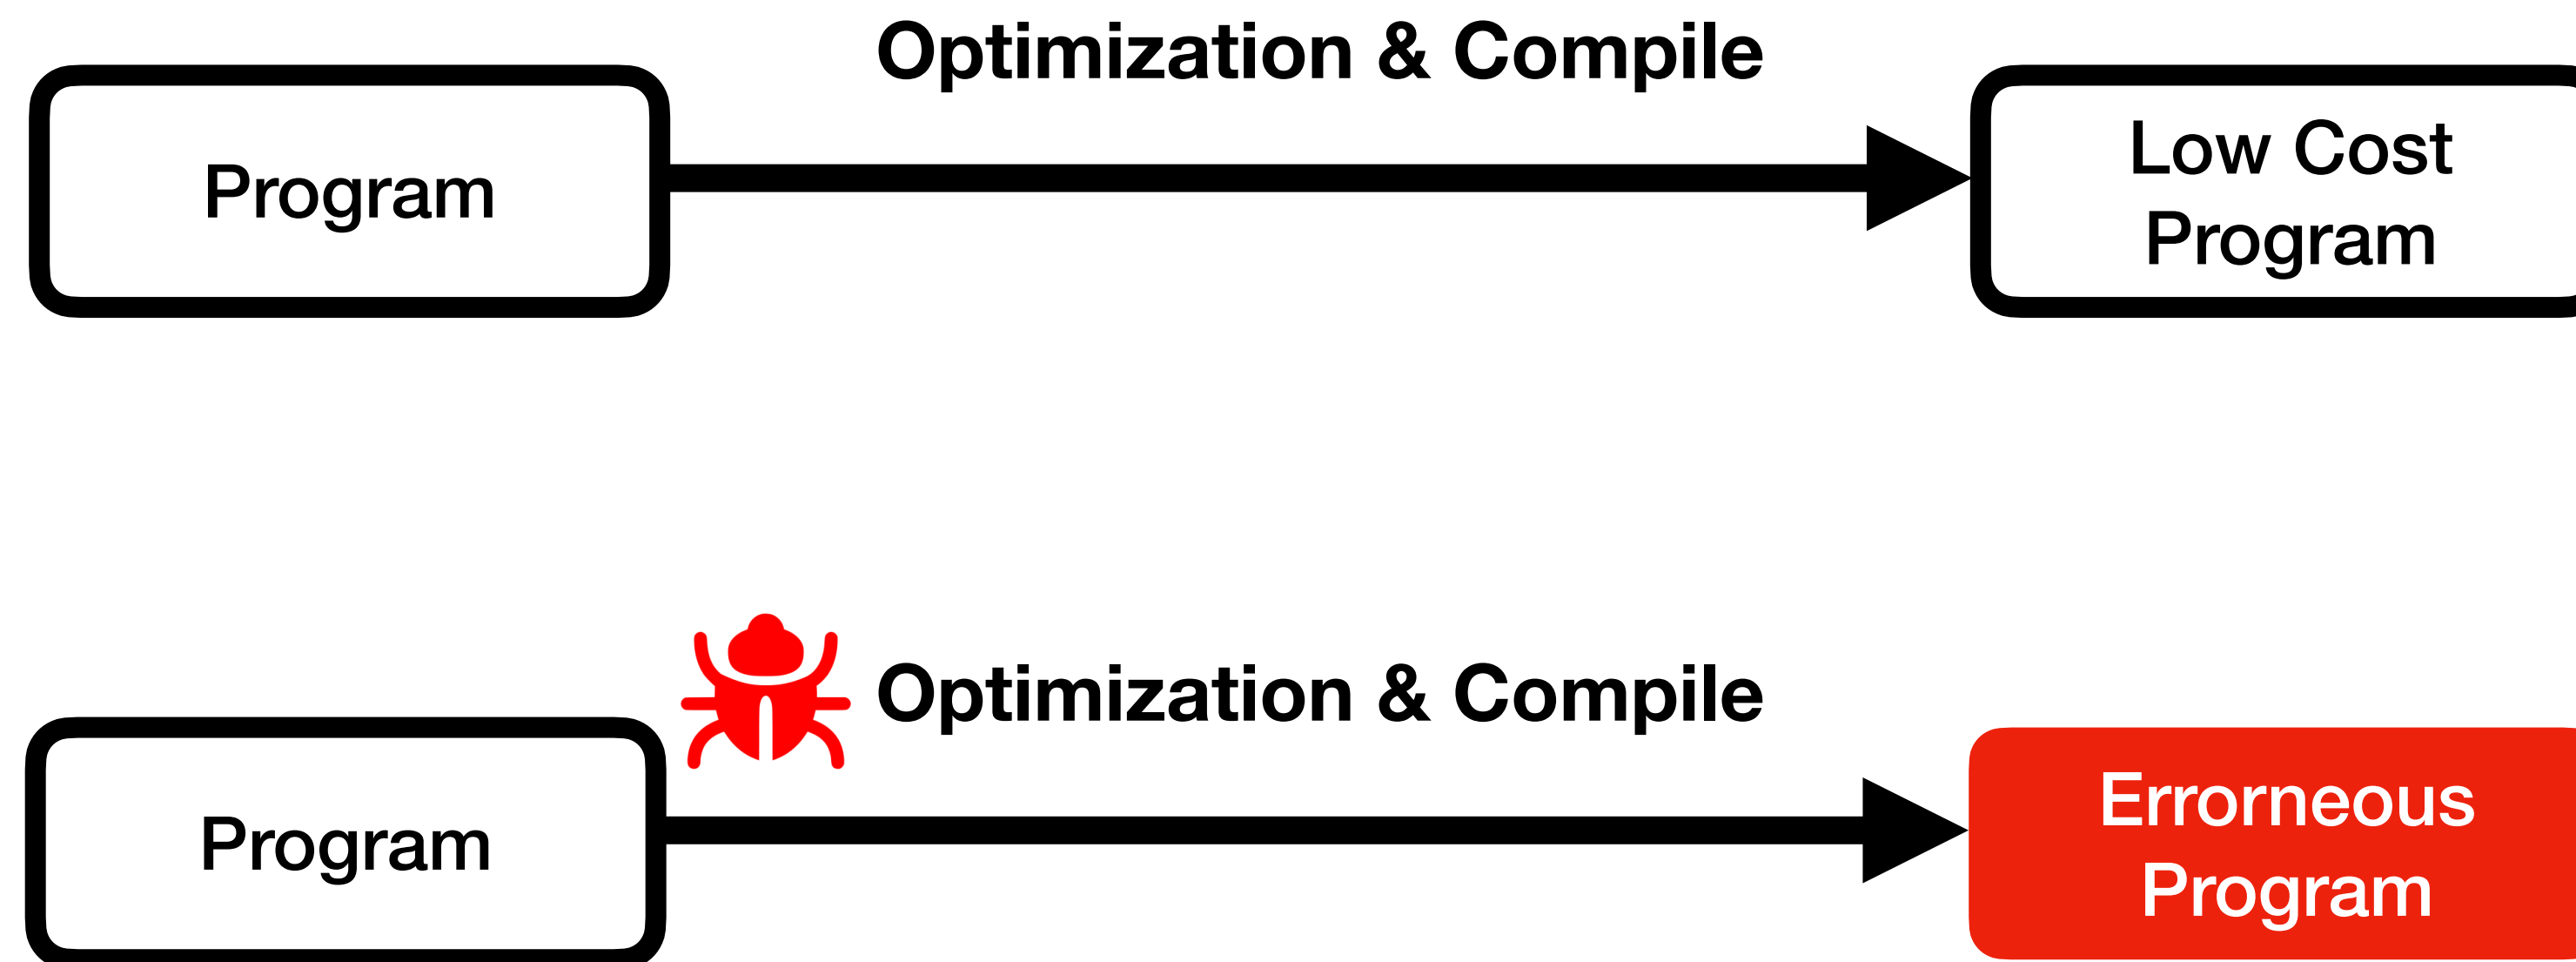
Low-cost Program

Motivation: Optimization Bugs



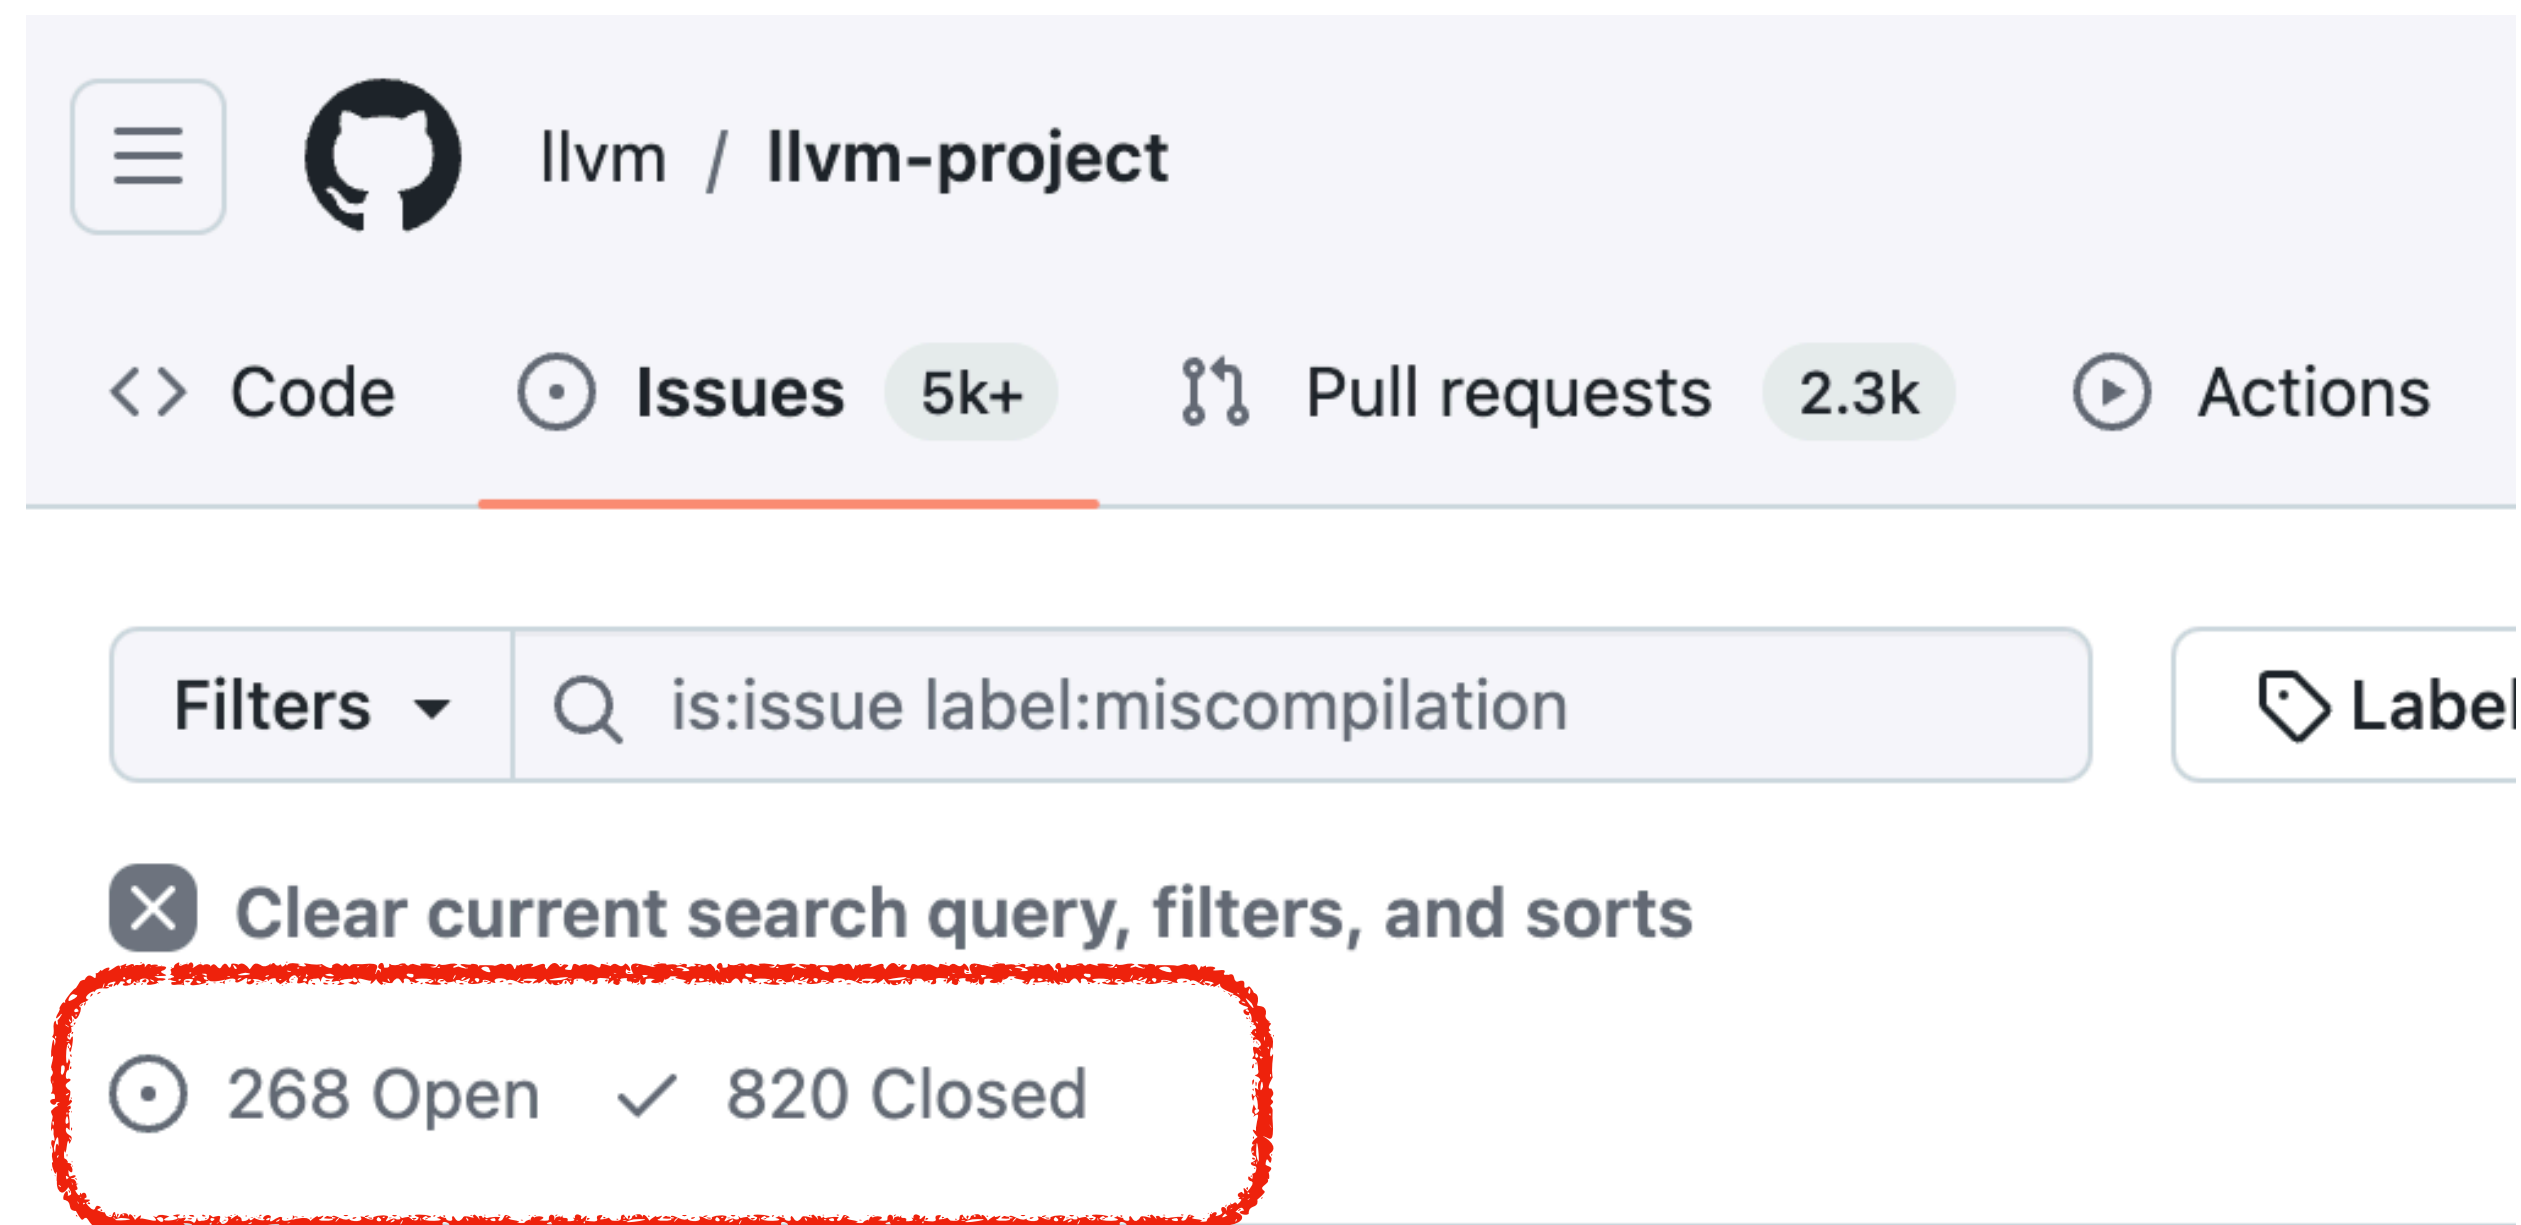
Motivation: Optimization Bugs

- Incorrect optimization can produce erroneous programs



LLVM's Bugs

- More than 1k bugs are reported



The screenshot shows the GitHub interface for the LLVM project. The top navigation bar includes the LLVM logo, the repository name 'llvm / llvm-project', and tabs for 'Code', 'Issues' (5k+), 'Pull requests' (2.3k), and 'Actions'. The 'Issues' tab is selected. Below the navigation bar, there is a search bar with the query 'is:issue label:miscompilation' and a 'Filters' dropdown. A button labeled 'Clear current search query, filters, and sorts' is visible. At the bottom, a red box highlights the issue status summary: '268 Open' and '820 Closed'.

llvm / llvm-project

<> Code Issues 5k+ Pull requests 2.3k Actions

Filters is:issue label:miscompilation Label

Clear current search query, filters, and sorts

268 Open ✓ 820 Closed

Challenge: LLVM's Complexity




- More than **2.6 M** lines of code
- 3,525 commits in a month (2024.03)
- 56 Contributors (Developers) in a month

Hard to verify compiler in real-world!

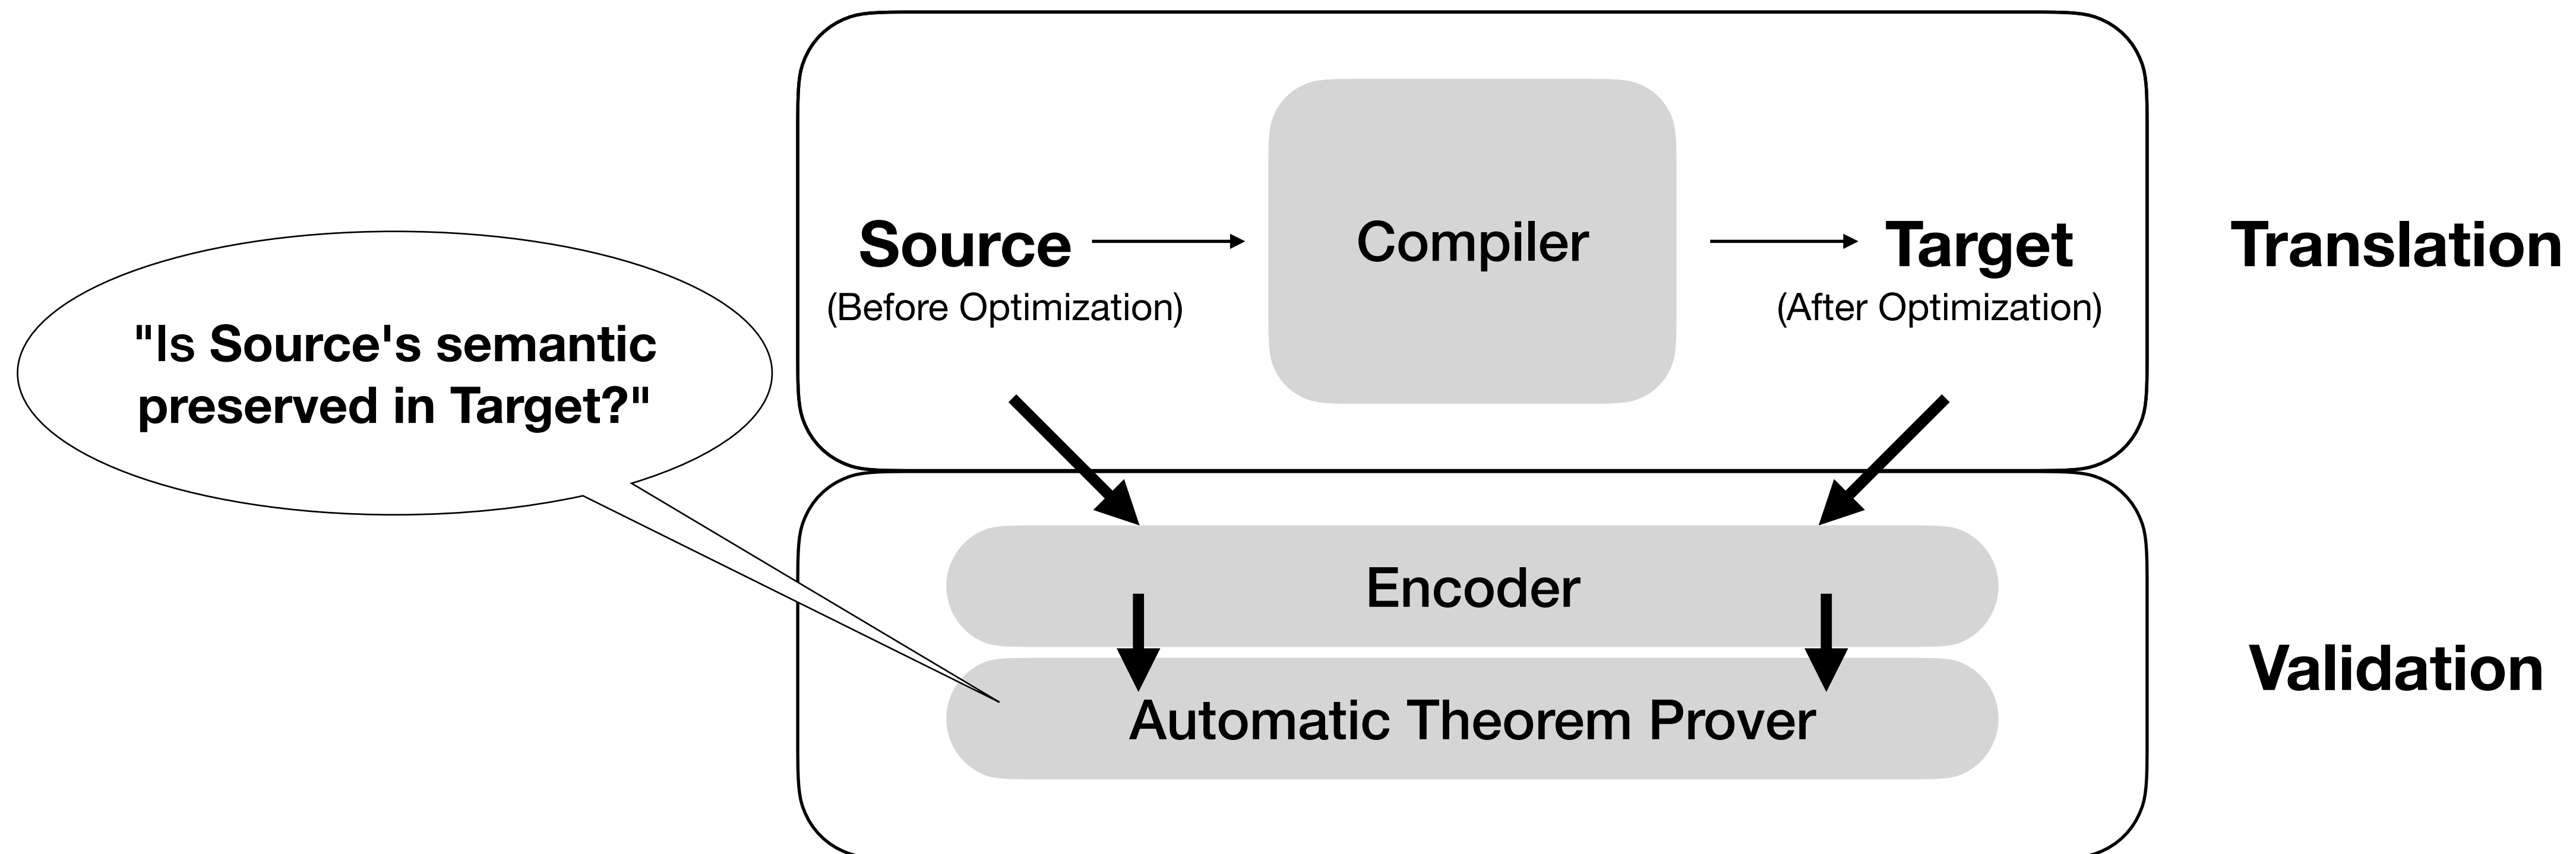
Idea: Translation Validation (TV)

- Check if translation preserves the program's semantics
- Agnostic to compiler's implementation, focusing only on language's semantics



Idea: Translation Validation (TV)

- Check if translation preserves the program's semantics
- Agnostic to compiler's implementation, focusing only on language's semantics



Example: How Alive2 Works?

Source

```
int foo(int X, int Y) {  
  int result = X + Y;  
  return result;  
}
```



Optimize



Target

```
int foo(int X, int Y) {  
  int result = X | Y;  
  return result;  
}
```

Example: How Alive2 Works?

Source

```
int foo(int X, int Y) {  
  int result = X + Y;  
  return result;  
}
```



Optimize



Target

```
int foo(int X, int Y) {  
  int result = X | Y;  
  return result;  
}
```

Encode : $\forall X \forall Y. X + Y$

Example: How Alive2 Works?

Source

```
int foo(int X, int Y) {  
  int result = X + Y;  
  return result;  
}
```



Optimize



Target

```
int foo(int X, int Y) {  
  int result = X | Y;  
  return result;  
}
```

Encode : $\forall X \forall Y. X + Y$

$\forall X \forall Y. X | Y$

Example: How Alive2 Works?

Source

```
int foo(int X, int Y) {  
  int result = X + Y;  
  return result;  
}
```



Optimize



Target

```
int foo(int X, int Y) {  
  int result = X | Y;  
  return result;  
}
```

Encode : $\forall X \forall Y. X + Y$

$\forall X \forall Y. X | Y$

Solve (SMT-Solver) : $\exists X \exists Y. X + Y \neq X | Y$ is satisfiable?

Example: How Alive2 Works?

Source

```
int foo(int X, int Y) {  
  int result = X + Y;  
  return result;  
}
```



Optimize



Target

```
int foo(int X, int Y) {  
  int result = X | Y;  
  return result;  
}
```

Encode : $\forall X \forall Y. X + Y$

$\forall X \forall Y. X | Y$

Solve (SMT-Solver) : $\exists X \exists Y. X + Y \neq X | Y$ is **satisfiable?**

Result : Satisfiable, Model: **X: 4, Y: 2**

Example: How Alive2 Works?

Source

```
int foo(int X, int Y) {  
  int result = X + Y;  
  return result;  
}
```



Target

```
int foo(int X, int Y) {  
  int result = X | Y;  
  return result;  
}
```

Encode : $\forall X \forall Y. X + Y$

$\forall X \forall Y. X | Y$

Solve (SMT-Solver) : $\exists X \exists Y. X + Y \neq X | Y$ is **satisfiable?**

Result : Satisfiable, Model: **X: 4, Y: 2**

Wrong Optimization! Semantic are not equivalent When **X** is 4 & **Y** is 2

How Is It Possible?


- LLVM has specification*
 - It is easy to encode semantic in a program

How Is It Possible?

- LLVM has specification*
 - It is easy to encode semantic in a program
- All programming languages are represented by LLVM IR, then optimized
 - Only need to encode the LLVM IR program

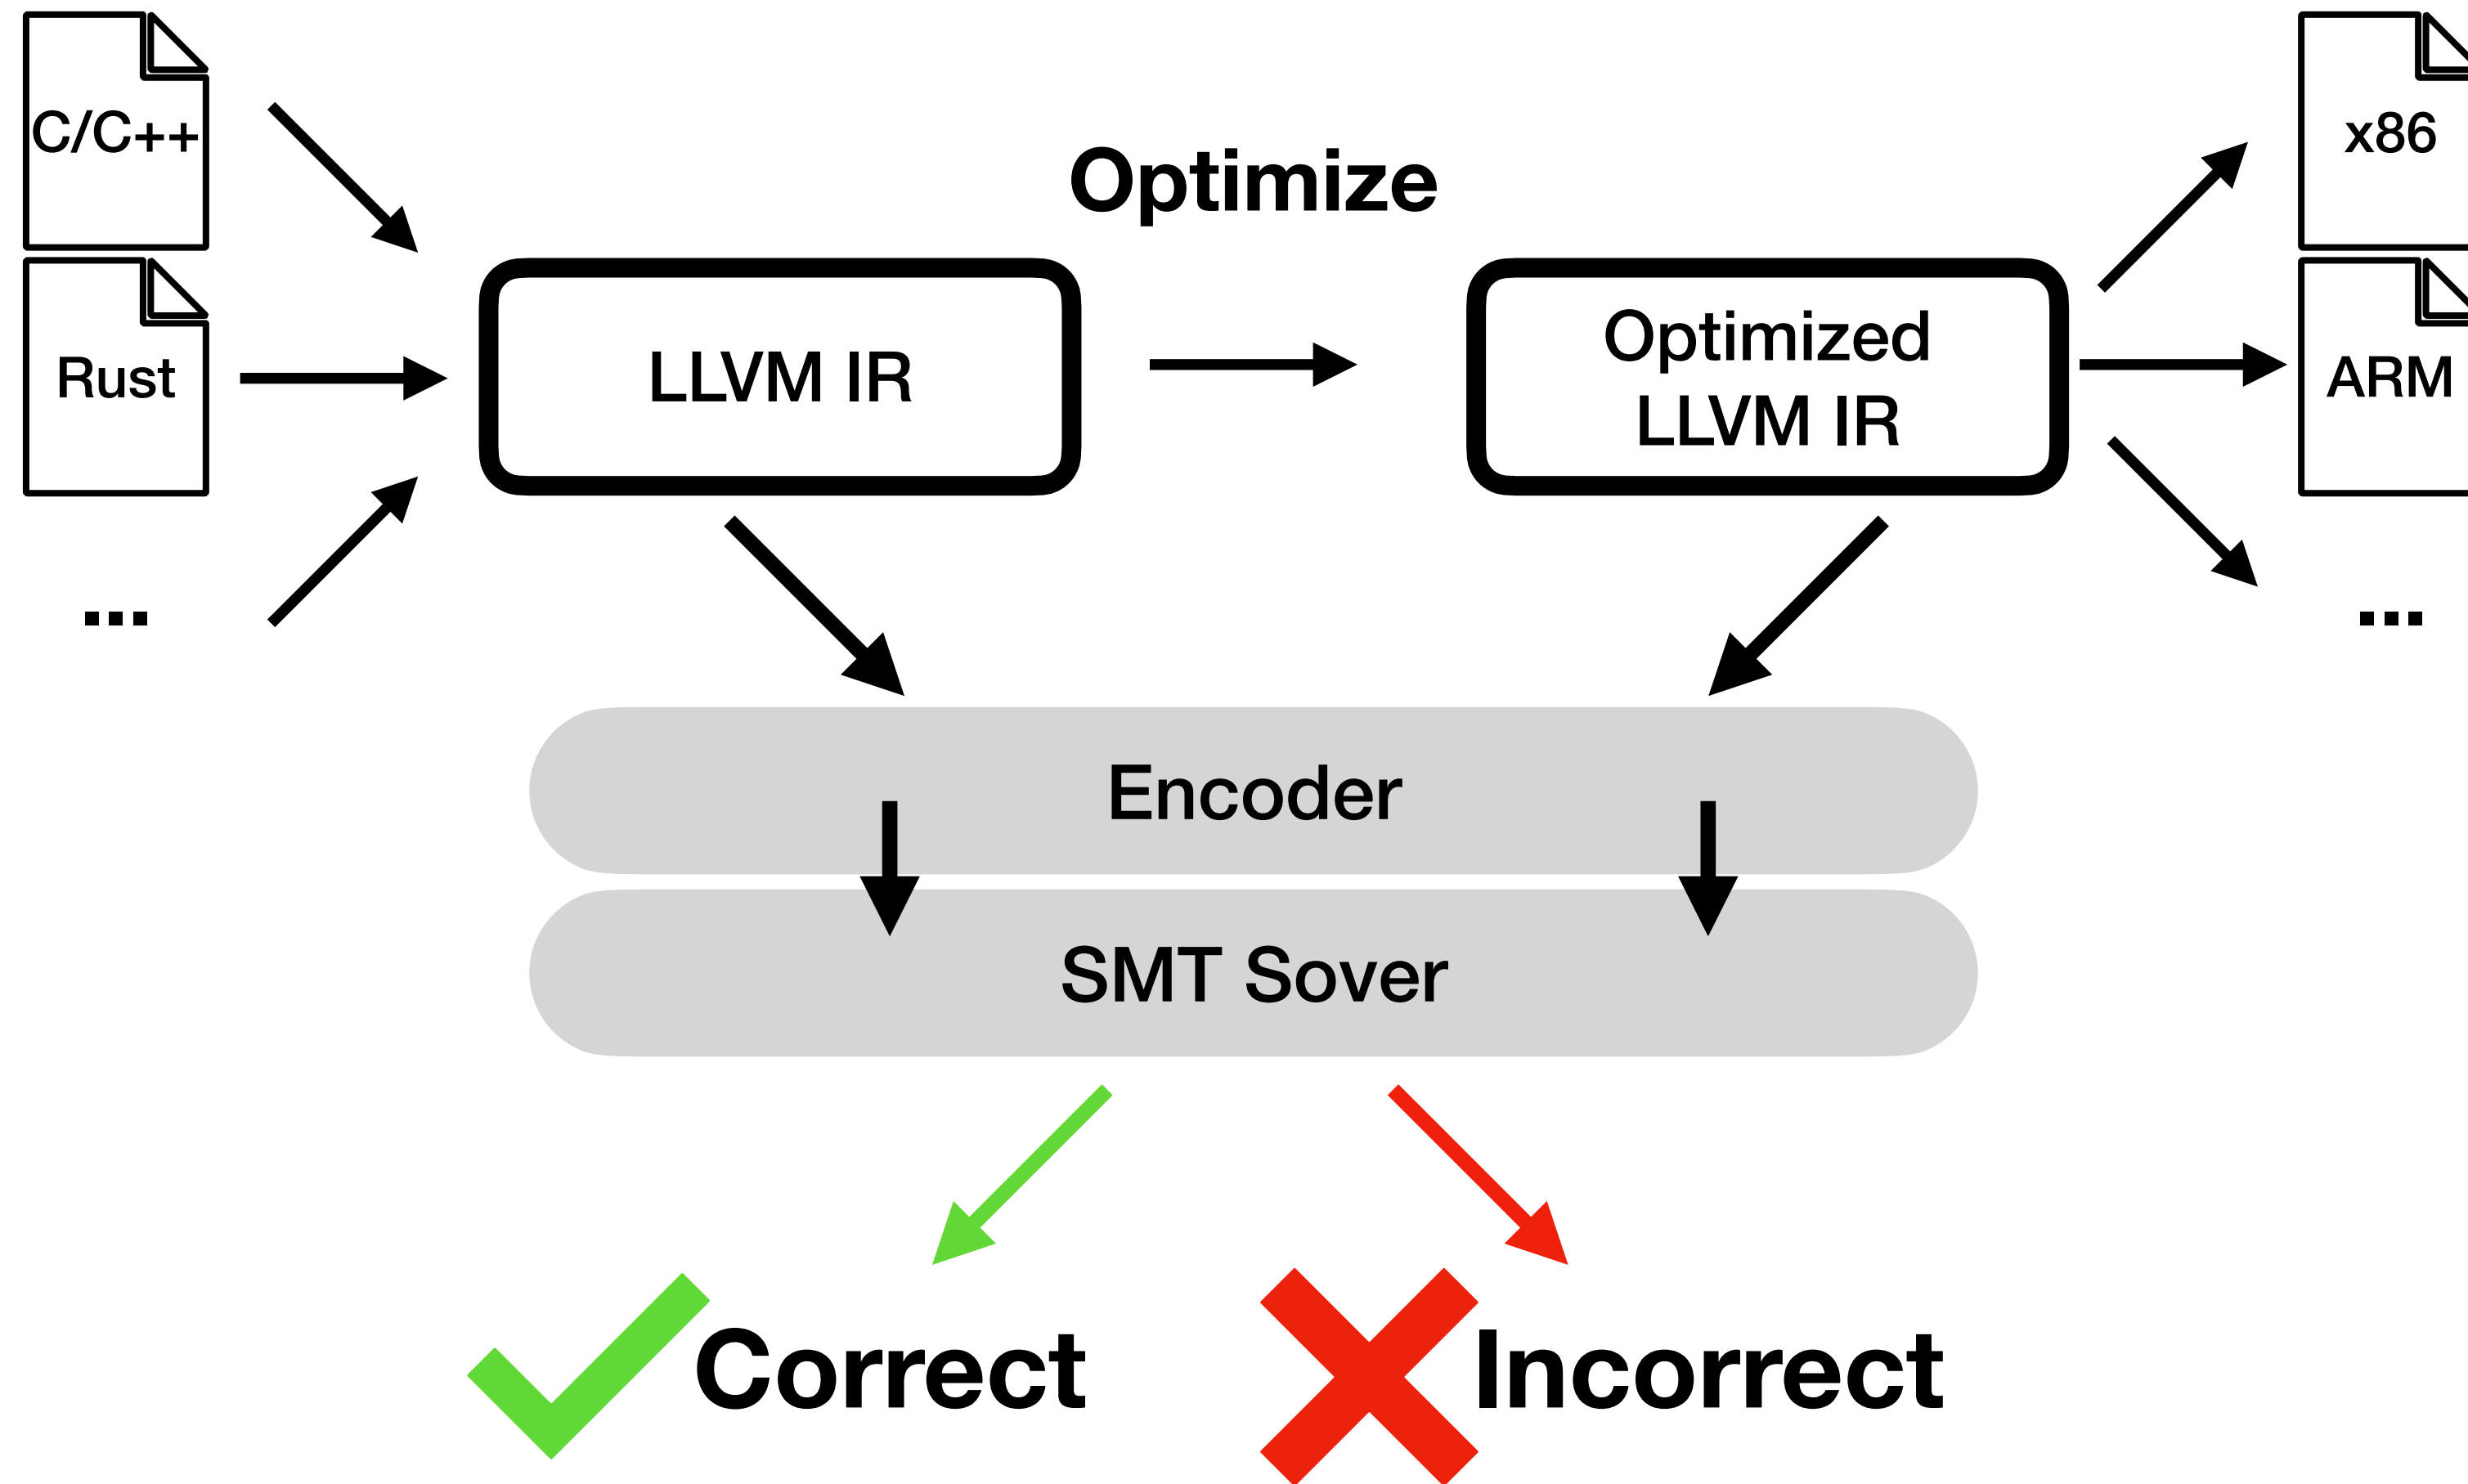
How Is It Possible?

- LLVM has specification*
 - It is easy to encode semantic in a program
- All programming languages are represented by LLVM IR, then optimized
 - Only need to encode the LLVM IR program



*<https://llvm.org/docs/LangRef.html>

Alive2




Appealing Result



47 bugs in LLVM Test Suits

Appealing Result




47 bugs in LLVM Test Suites



8 patches in LLVM Specification

Appealing Result



47 bugs in LLVM Test Suites



8 patches in LLVM Specification



Only **<3%** false positives

Technical Details

- **2 key technical details**
- What is correct optimization in LLVM?
 - Define and Check refinement of LLVM IR program
- How to improve precision of Alive2?
 - Unroll loops up to some bound

Technical Details

- 2 key technical details
- What is correct optimization in LLVM?
 - Define and Check refinement of LLVM IR program
- How to improve precision of Alive2?
 - Unroll loops up to some bound

Technical Details

- 2 key technical details
- What is correct optimization in LLVM?
 - Define and Check refinement of LLVM IR program
- How to improve precision of Alive2?
 - Unroll loops in program up to some bound

LLVM IR's Refinement

- **Undefined Behavior** is important
 - division by zero, shift past bitwidth ...

LLVM IR's Refinement

- **Undefined Behavior** is important
 - division by zero, shift past bitwidth ...

// division by zero

```
int foo(int X) {  
    int result = X / 0;  
    return result;  
}
```



Undef

(... -2, -1, 0, 1, 2 ...)

// shift past bitwidth

```
int foo(int X) {  
    int result = X << 100;  
    return result;  
}
```



Undef

(... -2, -1, 0, 1, 2 ...)

LLVM IR's Refinement

- **Undefined Behavior** is important

```
int foo(int X) {  
    int result = (X / 0) + 1;  
    return result;  
}
```

LLVM IR's Refinement

- **Undefined Behavior** is important

```
int foo(int X) {  
  int result = (X / 0) + 1;  
  return result;  
}
```

→ **Undef + 1**
(... -2, -1, 0, 1, 2 ...)

LLVM IR's Refinement

- **Undefined Behavior** is important

```
int foo(int X) {  
  int result = (X / 0) + 1;  
  return result;  
}
```

→ **Undef + 1**
(... -2, -1, 0, 1, 2 ...)

→ **Undef**
(... -2, -1, 0, 1, 2 ...)


LLVM IR's Refinement

- **Undefined Behavior** is important




Correct Optimization!

```
int foo(int X) {  
  int result = (X / 0) + 1;  
  return result;  
}
```



Undef + 1

(... -2, -1, 0, 1, 2 ...)



Undef

(... -2, -1, 0, 1, 2 ...)



```
int foo(int X) {  
  return 1;  
}
```


LLVM IR's Refinement

- **Undefined Behavior** is important




Correct Optimization!

```
int foo(int X) {  
  int result = (X / 0) + 1;  
  return result;  
}
```



Undef + 1

(... -2, -1, 0, 1, 2 ...)



Undef

(... -2, -1, 0, 1, 2 ...)



```
int foo(int X) {  
  return 1;  
}
```

```
int foo(int X) {  
  int result = (X / 0) | 1;  
  return result;  
}
```


LLVM IR's Refinement

- **Undefined Behavior** is important




Correct Optimization!

```
int foo(int X) {  
  int result = (X / 0) + 1;  
  return result;  
}
```



Undef + 1

(... -2, -1, 0, 1, 2 ...)




Undef

(... -2, -1, 0, 1, 2 ...)



```
int foo(int X) {  
  return 1;  
}
```

```
int foo(int X) {  
  int result = (X / 0) | 1;  
  return result;  
}
```



Undef | 1

(... -3, -1, 1, 3 ...)

LLVM IR's Refinement

- **Undefined Behavior** is important



Correct Optimization!

```
int foo(int X) {  
  int result = (X / 0) + 1;  
  return result;  
}
```

→ **Undef + 1**
(... -2, -1, 0, 1, 2 ...)

→ **Undef**
(... -2, -1, 0, 1, 2 ...)

```
int foo(int X) {  
  return 1;  
}
```

```
int foo(int X) {  
  int result = (X / 0) | 1;  
  return result;  
}
```

→ **Undef | 1**
(... -3, -1, 1, 3 ...)

→ **Undef**
(... -2, -1, 0, 1, 2 ...)

```
int foo(int X) {  
  return 4;  
}
```

LLVM IR's Refinement

- **Undefined Behavior** is important




Correct Optimization!

```
int foo(int X) {  
  int result = (X / 0) + 1;  
  return result;  
}
```

→ **Undef + 1**
(... -2, -1, 0, 1, 2 ...)

→ **Undef**
(... -2, -1, 0, 1, 2 ...)

```
int foo(int X) {  
  return 1;  
}
```



InCorrect Optimization!

```
int foo(int X) {  
  int result = (X / 0) | 1;  
  return result;  
}
```

→ **Undef | 1**
(... -3, -1, 1, 3 ...)

→ **Undef**
(... -2, -1, 0, 1, 2 ...)

```
int foo(int X) {  
  return 4;  
}
```

LLVM IR's Refinement

- **Alive2** defined every case of LLVM IR's refinement


$$R[\%a] = (\text{ite}(\text{isundef}_{\%a}, \text{undef}_1, \%a), \text{ispoison}_{\%a})$$
$$R[\%b] = (\text{ite}(\text{isundef}_{\%b}, \text{undef}_2, \%b), \text{ispoison}_{\%b})$$
$$R[\%t] = (\text{ite}(\text{isundef}_{\%a}, \text{undef}_3, \%a) + \text{ite}(\text{isundef}_{\%a}, \text{undef}_4, \%a), \\ \text{ispoison}_{\%a})$$
$$R[\%c] = (\text{ite}(\text{ite}(\text{isundef}_{\%a}, \text{undef}_5, \%a) + \\ \text{ite}(\text{isundef}_{\%a}, \text{undef}_6, \%a) = 0, 1, 0), \text{ispoison}_{\%a})$$
$$R[\%q] = (\text{shl}(\%a, 2), \text{false})$$
$$R[\%r] = (\text{ite}(\text{isundef}_{\%b}, \text{undef}_7, \%b) \& 1, \text{ispoison}_{\%b})$$

...

Alive2's Loop Unrolling

- Alive2 wants "**Bug means a bug**"
- Alive2 Unroll loops in program up to some bound

```
int foo(int X) {  
  int j = 3;  
  for (int i = 0; i < 3; i++)  
  {  
    j--;  
  }  
  return X + j;  // X + 0
```




```
int foo(int X) {  
  return X;  // X + 0
```

Alive2's Loop Unrolling

- Alive2 wants "**Bug means a bug**"
- Alive2 Unroll loops in program up to some bound

```
int foo(int X) {  
  int j = 3;  
  for (int i = 0; i < 3; i++)  
  {  
    j--;  
  }  
  return X + j;  // X + 0
```



```
int foo(int X) {  
  int j = 3;  
  j--;  
  j--;  
  j--;  
  return X + j;  // X + 0
```

// unrolling



```
int foo(int X) {  
  return X;  // X + 0
```


Evaluation: Precision

- Translation Validation of LLVM's Unit Tests

All TestSuits	Supported	True Alarms	Reported
168,000	36,000	121	43

Evaluation: Scalability

- Translation Validation for Applications

Programs	Total	True Alarms	False Alarms	Failed	Unsupported
bzip2	2.2K	333	10	735	1,125
gzip	2.6K	884	4	965	754
oggenc	1.8K	440	4	660	663
ph7	5.6K	1,393	28	1,372	2,755
SQLite3	12.2K	2,314	38	2,202	7,543

Evaluation: Scalability

- Translation Validation for Applications

Only <3% False Positives!

Programs	Total	True Alarms	False Alarms	Failed	Unsupported
bzip2	2.2K	333	10	735	1,125
gzip	2.6K	884	4	965	754
oggenc	1.8K	440	4	660	663
ph7	5.6K	1,393	28	1,372	2,755
SQLite3	12.2K	2,314	38	2,202	7,543

Evaluation

- **Updates to the LLVM IR Semantics**
 - In encoding LLVM IR, Ambiguously written specifications were founded
 - After discussions with LLVM communities, 8 patches applied to spec

Conclusion

- Translation validator for LLVM middle-end optimization
 - TV is suitable for LLVM's optimization
- Formally defined refinement & Precise validation
 - Alive2 is suitable for real-world validation
- Found new bugs in a LLVM
 - 47 bugs in LLVM Test Suits
 - 8 patches in LLVM Specification

Review (My Opinion)

- Positive aspects
 - A useful testing system to find optimization bugs in LLVM
 - Fast, accurate, and able to identify the reasons for bugs
- Negative aspects
 - When new features are added to LLVM IR, new encodings need to be added
 - Since encoding is done manually by humans, problems can arise (Alive2's bug)